

Page 1 of a long scanned report.

Retrieval-augmented generation pipelines depend on clean, well-segmented text. When a source document is a scan, no native text layer exists, so optical understanding becomes the only path to usable content for downstream embedding.

Page 2 of a long scanned report.

Retrieval-augmented generation pipelines depend on clean, well-segmented text. When a source document is a scan, no native text layer exists, so optical understanding becomes the only path to usable content for downstream embedding.

Page 3 of a long scanned report.

Retrieval-augmented generation pipelines depend on clean, well-segmented text. When a source document is a scan, no native text layer exists, so optical understanding becomes the only path to usable content for downstream embedding.

Page 4 of a long scanned report.

Retrieval-augmented generation pipelines depend on clean, well-segmented text. When a source document is a scan, no native text layer exists, so optical understanding becomes the only path to usable content for downstream embedding.

Page 5 of a long scanned report.

Retrieval-augmented generation pipelines depend on clean, well-segmented text. When a source document is a scan, no native text layer exists, so optical understanding becomes the only path to usable content for downstream embedding.

Page 6 of a long scanned report.

Retrieval-augmented generation pipelines depend on clean, well-segmented text. When a source document is a scan, no native text layer exists, so optical understanding becomes the only path to usable content for downstream embedding.

Page 7 of a long scanned report.

Retrieval-augmented generation pipelines depend on clean, well-segmented text. When a source document is a scan, no native text layer exists, so optical understanding becomes the only path to usable content for downstream embedding.

Page 8 of a long scanned report.

Retrieval-augmented generation pipelines depend on clean, well-segmented text. When a source document is a scan, no native text layer exists, so optical understanding becomes the only path to usable content for downstream embedding.

Page 9 of a long scanned report.

Retrieval-augmented generation pipelines depend on clean, well-segmented text. When a source document is a scan, no native text layer exists, so optical understanding becomes the only path to usable content for downstream embedding.

Page 10 of a long scanned report.

Retrieval-augmented generation pipelines depend on clean, well-segmented text. When a source document is a scan, no native text layer exists, so optical understanding becomes the only path to usable content for downstream embedding.

Page 11 of a long scanned report.

Retrieval-augmented generation pipelines depend on clean, well-segmented text. When a source document is a scan, no native text layer exists, so optical understanding becomes the only path to usable content for downstream embedding.

Page 12 of a long scanned report.

Retrieval-augmented generation pipelines depend on clean, well-segmented text. When a source document is a scan, no native text layer exists, so optical understanding becomes the only path to usable content for downstream embedding.

Page 13 of a long scanned report.

Retrieval-augmented generation pipelines depend on clean, well-segmented text. When a source document is a scan, no native text layer exists, so optical understanding becomes the only path to usable content for downstream embedding.

Page 14 of a long scanned report.

Retrieval-augmented generation pipelines depend on clean, well-segmented text. When a source document is a scan, no native text layer exists, so optical understanding becomes the only path to usable content for downstream embedding.

Page 15 of a long scanned report.

Retrieval-augmented generation pipelines depend on clean, well-segmented text. When a source document is a scan, no native text layer exists, so optical understanding becomes the only path to usable content for downstream embedding.

Page 16 of a long scanned report.

Retrieval-augmented generation pipelines depend on clean, well-segmented text. When a source document is a scan, no native text layer exists, so optical understanding becomes the only path to usable content for downstream embedding.

Page 17 of a long scanned report.

Retrieval-augmented generation pipelines depend on clean, well-segmented text. When a source document is a scan, no native text layer exists, so optical understanding becomes the only path to usable content for downstream embedding.

Page 18 of a long scanned report.

Retrieval-augmented generation pipelines depend on clean, well-segmented text. When a source document is a scan, no native text layer exists, so optical understanding becomes the only path to usable content for downstream embedding.

Page 19 of a long scanned report.

Retrieval-augmented generation pipelines depend on clean, well-segmented text. When a source document is a scan, no native text layer exists, so optical understanding becomes the only path to usable content for downstream embedding.

Page 20 of a long scanned report.

Retrieval-augmented generation pipelines depend on clean, well-segmented text. When a source document is a scan, no native text layer exists, so optical understanding becomes the only path to usable content for downstream embedding.

Page 21 of a long scanned report.

Retrieval-augmented generation pipelines depend on clean, well-segmented text. When a source document is a scan, no native text layer exists, so optical understanding becomes the only path to usable content for downstream embedding.

Page 22 of a long scanned report.

Retrieval-augmented generation pipelines depend on clean, well-segmented text. When a source document is a scan, no native text layer exists, so optical understanding becomes the only path to usable content for downstream embedding.

Page 23 of a long scanned report.

Retrieval-augmented generation pipelines depend on clean, well-segmented text. When a source document is a scan, no native text layer exists, so optical understanding becomes the only path to usable content for downstream embedding.

Page 24 of a long scanned report.

Retrieval-augmented generation pipelines depend on clean, well-segmented text. When a source document is a scan, no native text layer exists, so optical understanding becomes the only path to usable content for downstream embedding.

Page 25 of a long scanned report.

Retrieval-augmented generation pipelines depend on clean, well-segmented text. When a source document is a scan, no native text layer exists, so optical understanding becomes the only path to usable content for downstream embedding.