

Retrieval-augmented generation pipelines depend on clean, well-segmented text. When a source document is a scan, no native text layer exists, so optical understanding becomes the only path to usable content for downstream embedding.

This page evaluates whether the vision fallback recovers readable paragraphs from a rasterized image. The extracted text should preserve reading order and omit running headers and footers that do not belong to the body.